

Introduction

Performance model of matrix multiplication architecture. For now, the functional unit computes the outer product GEMM kernel

$$C_{M,N} = A_{M,K} * B_{K,N} = \sum_i A_{M,i} \times B_{i,N}.$$

We would like to characterize how architectural parameters such as number and size of matrix registers, supported datatype, scalability of register dimensions, instructions, etc, affect performance metrics such as kernel latency, throughput (operations per cycle), and required bandwidth of memory and matrix register file, and cache size to maximize throughput.

We also explore the sensitivity of these performance metrics to software parameters such as matrix dimensions.

In the future we plan to use RTL simulation to see how these performance metrics are reduced when cache size and memory bandwidth to not meet these requirements.

Functional Unit Layout

A a two dimensional grid of registers and multiply-accumulate (macc) units, stores partial products of a subset, or tile, of the output matrix, $C_i = A_{M,i} \times B_{i,N}$. A mesh of wires connects the macc units to columns of A and rows of B. This outer-product-centric approach to accelerating GEMM kernels encourages C to be held in accumulation registers distributed with the MACC units.

	B ₀	B ₁	...	B _{ml}
A ₀	C ₀₀ += A ₀ * B ₀	C ₀₁ += A ₀ * B ₁	...	C _{0,ml} += A ₀ * B _{ml}
A ₁	C ₁₀ += A ₁ * B ₀	C ₁₁ += A ₁ * B ₁	...	C _{1,ml} += A ₁ * B _{ml}
...
A _{vl}	C _{vl,0} += A _{vl} * B ₀	C _{vl,1} += A _{vl} * B ₁	...	C _{vl,ml} += A _{vl} * B _{ml}

Figure 1: image

Fig 1. Physical Layout of Outer Product

The outer product unit (OPU) can execute NM MACC operations per cycle, and requires $N + M$ elements loaded from memory. For a machine with vector length $vl = N = M$, the OPU performance scales as $NM = vl^2$ MACCs per cycle, so performance scales quadratically with the vector length.

The architecture provides loading tiles of A, B and C, and performing outer-product-accumulate on C.

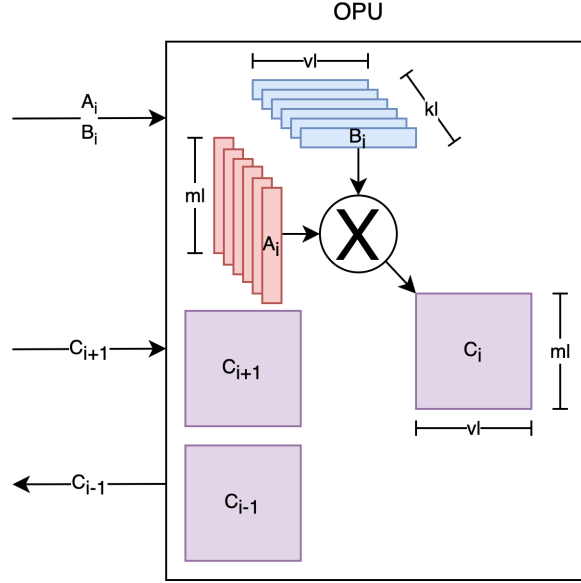


Figure 2: image

GEMM Outer Product Kernel

Linear algebra libraries decompose large GEMMs into micro-kernels, such as the output stationary outer product kernel below:

Decomposition of GEMM into Outer Product Micro-Kernels [1]. where the register and cache tile dimensions n_r , m_r , k_c correspond to our model parameters like so: $n_r = v_l$, $m_r = m_l$

The micro kernel at the bottom of the above figure can be written as follows:

```

vset vl N
mset ml M
mld mrc C

ukernel_x_kc:
    vld vra A[:,k]
    vld vrb B[k,:]
    opacc mrc vra vrb
    ... iterate (kc)x
    beq ... ukernel
mst mrc C

```




Figure 3: image

where `mld` and `mst` move matrix tiles into and out of the matrix registers within the outer product unit (note `mld` and `mst` need the unit-stride dimension of the matrix in memory to compute the stride between vector rows). A scalable dimension `ml` has been added to allow the number of vector rows held in the OPU to vary across machines.

The micro-kernel is computed by calling the `opacc` instruction k_c times, as shown in the above Fig. The k_c iterations of micro-kernel can be reduced by introducing the parameter k_l , which is the number of accumulates per outer product instruction.

```

vset vl N
mset ml M
mld mrc C

ukernel_x_(kc/kl):
    kset kl kc
    mld mra A
    mld mrb B
    opacc mrc mra mrb
    ... iterate (kc/kl)x
    beq ... ukernel
mst mrc C

```

Now A and B are tiles with inner dimension k_l , instead of vectors.

In the inner loop of this output stationary dataflow, $N + M$ vector elements are loaded and NM MACC operations are performed. For a machine with vector length $vl = N = M$, the fraction of compute operations to memory accesses, or operational intensity, increases with the vector length of the machine, $OI = \frac{NM}{N+M} = vl$.

Scaling vl is limited in part by available register and cache capacity required to stage the next set of outer products such that the functional unit is fully utilized.

GEMM Dataflows

First, consider GEMM with small matrix dimensions, where A, B, and C all fit in the cache. Tiles of A, B, and C are held in matrix registers. The matrix r

To hide memory latency t_{ld} , p_{mem} threads must send memory load requests in parallel, requiring p_{mem} matrix tile registers to receive the load data. A cache can be used to buffer data loaded for the next batch of GEMMs and data to be stored from the last batch of GEMMs.

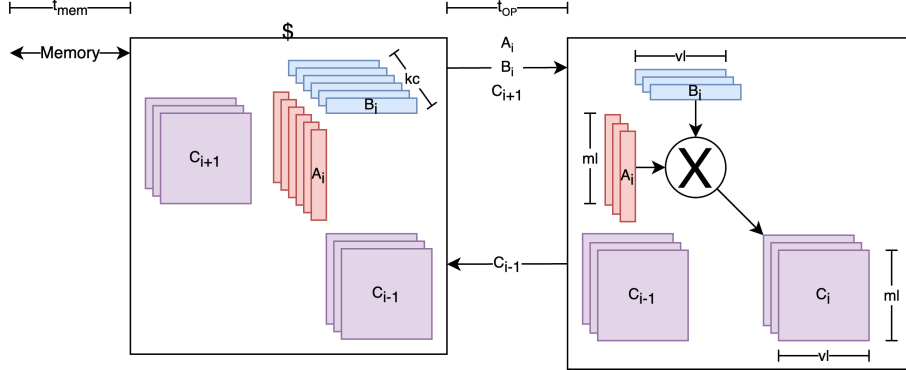


Figure 4: image

If the C matrix is small enough to fit in an matrix register, then the parallel memory requests must come from independent threads, and the only opportunity for memory data reuse is in the K-dimension, because tile C will be reused K times.

Operation Latency

The latency of an outer product micro-kernel, $t_{\mu k}$ depends on micro architecture decisions such as number of read and write ports between the cache and matrix register file.

Assuming there is sufficient buffering and communication bandwidth to overlap

the outer product computation with loading and storing tiles of C, the minimum possible latency is

$$t_{\mu k} = k_c$$

which requires the load bandwidth between the cache and matrix register file (mrf) to be,


$$BW_{mrf} = \frac{k_c(m_l + v_l) + m_l v_l}{k_c}$$

to load the A and B tiles for current OPACC as well as the C tile for the next micro-kernel. It is also possible to reduce the required bandwidth by reducing the macc operations per cycle and thus increasing the operation latency.

Future work is to build a matrix unit generator in RTL to find cycle-accurate latency over many micro-architecture parameters.

Memory Latency

Here is an example of a simple single-threaded schedule:



To keep the OPU fully utilized we need to request overlapping loads from memory. The memory request for the next A B C tiles must complete by $t = t_{ld} + m_l + k$, thus a new memory request must be issued every k cycles. To provide register storage for all in-flight loads, we need a number of parallel overlapping loads equal to:

$$p_{mem} = \frac{t_{ld} + m_l}{t_{\mu k}} = \frac{t_{ld} + m_l}{k_c}$$

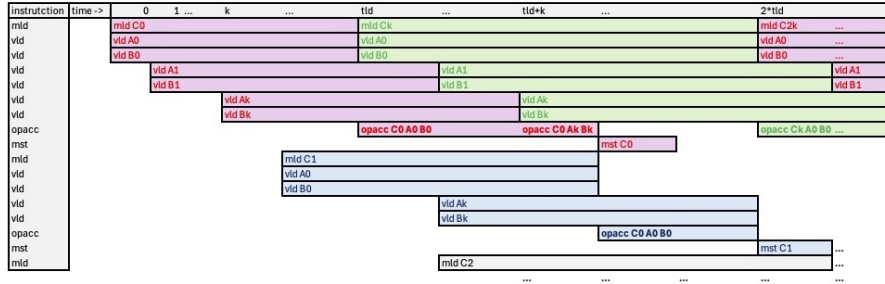


Figure 5: image

In this model we consider a range of memory latencies from local caches (1-5 cycles) to shared cache (20-50 cycles) to DRAM memory (100-300 cycles).

Worst case scenario, $p_{mem} \approx \frac{200}{4} = 50$. To avoid implementing 50 matrix registers, data should be prefetched into a smaller cache. By prefetching into the L2 cache, number of registers for full throughput is roughly $p_{mem} \approx \frac{40}{4} = 10$. L2 cache can be designed to have lower latency, allowing for even fewer registers.

If there are not enough registers to maintain full throughput for a given memory latency, utilization drops and can be expressed as:

$$U_{op} = \frac{t_{opacc}}{t_{uk}} = \max(1, \frac{p_{mem}k_l}{t_{mem}+k_l})$$

resulting in reduced operations per cycle, $U_{op}m_lv_l$

Cache Capacity

Over one batch of p_{mem} GEMMs, required cache capacity to fully utilize the OPU is

$$p_{mem}(2C_t + A_t + B_t) = p_{mem}(2m_lv_l + m_lk_c + k_cv_l) \text{ elements,}$$

where v_l , m_l , k_c , are the dimensions of tiles C_t , A_t , B_t . In general, operational intensity is maximized when matrix tiles are square. If $v_l = m_l$, required cache capacity can be expressed as

$$c_{L2} = p_{mem}(2m_l + 2k_c) \text{ vectors of length } v_l.$$

Memory Bandwidth

Over one batch of p_{mem} GEMMs, required memory bandwidth, is

$$BW_{mem} = \frac{c_{L2}}{t_{\mu k}} = \frac{m_l+k_c}{k_c} \text{ vector loads per cycle,}$$

Bandwidth for Large K Dimension

Now we see that as K is large, we achieve low bandwidth in addition to high operational intensity:

$$BW_{mem} \approx 1 \text{ vector load per cycle.}$$

In this ideal case, the cache simply needs to buffer the next A and B tiles that will be used to accumulate C, and the overhead of loading and storing C tiles is negligible.

Large M and N Dimensions

In general N and M are similar in size to K. If the dimensions N and M of matrices A B and C are sufficiently large, one tile of B can be reused for multiple C and A tiles, as shown in the first loop around the micro-kernel. This reduces the required cache capacity to

$$c_{L2} = k_c(p_{mem} + 1) \text{ vectors}$$

and memory bandwidth to

$$BW_{mem} = \frac{k_c(p_{mem}+1)}{k_c p_{mem}} = 1 + \frac{1}{p_{mem}} \text{ vectors per cycle.}$$

Bandwidth for Small K

GEMM with small K is more challenging because loading and storing C tiles begins to significantly increase $t_{\mu k}$ and thus degrade operational intensity.

To support applications with small K, the cache can buffer tiles C_{i-1} and C_{i+1} and move them in and out of matrix registers. The required cache capacity becomes

$$c_{L2} = 2p_{mem}(k_c + m_l) \text{ vectors}$$

and latency of the micro-kernel becomes

$$t_{\mu k} = p_{mem} \max(k_c, m_l)$$

Thus required memory bandwidth is

$$BW_{mem} = \frac{2(k_c+m_l)}{\max(m_l, k_c)} \text{ vector loads per cycle,}$$

Usage

We have discussed some design points in which the performance metrics have simple expressions. This model estimates the performance metrics over arbitrary matrix dimensions K, M, N, datatype, number of matrix registers, and instruction parameters such as v_l , m_l , and k_c . The model also includes dependence on micro-architectural parameters such as cache size, buffering, etc.

For example usage, see `dataflow_model.ipynb`

Reference

- [1] F. G. Van Zee and T. M. Smith, “Implementing high-performance complex matrix multiplication,” ACM Transactions on Mathematical Software, 2016